

Lab: Easy mode

Frank Edwards

4/24/2020

GLMs

- We can fit a frequentist regression model with a normal likelihood (using OLS) with the `lm()` function.

Refresher: lm()

- We can fit a frequentist regression model with a normal likelihood (using OLS) with the `lm()` function.
- R formulas take on an outcome \sim predictor1 + predictor2, data structure

```
library(broom)
data(WaffleDivorce)
m0<-lm(Divorce ~ Marriage, data = WaffleDivorce)
tidy(m0)
```

```
## # A tibble: 2 x 5
##   term          estimate std.error statistic   p.value
##   <chr>         <dbl>     <dbl>     <dbl>   <dbl>
## 1 (Intercept)    6.08      1.31      4.63 0.0000278
## 2 Marriage       0.179     0.0642    2.79 0.00751
```

Refresher: transformations and interactions

- We can use math functions like `log()`, `exp()` and `sqrt()` on outcomes or predictors
- To use other math, wrap a statement in `I()`

```
m1<-lm(log(Divorce) ~ sqrt(Marriage) + I(WaffleHouses / Population), data = WaffleDivorce)
tidy(m1)
```

```
## # A tibble: 3 x 5
##   term                estimate std.error statistic    p.value
##   <chr>                <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)          1.46      0.263      5.55 0.00000128
## 2 sqrt(Marriage)        0.170    0.0588      2.89 0.00586
## 3 I(WaffleHouses/Population) 0.00699  0.00272      2.57 0.0134
```

```
m2<-lm(Divorce ~ Marriage * South, data = WaffleDivorce)
tidy(m2)
```

```
## # A tibble: 4 x 5
##   term                estimate std.error statistic    p.value
##   <chr>                <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)          6.73      1.28      5.27 0.00000351
## 2 Marriage              0.129    0.0627      2.06 0.0448
## 3 South                -6.64      4.15     -1.60 0.117
## 4 Marriage:South         0.385    0.201      1.92 0.0613
```

The GLM function and model families

- We can estimate generalized linear models using `glm()`
- The family argument specifies the likelihood: family = “binomial”, “gaussian”, “Gamma”, “poisson” are most common
- The canonical link function for each family is specified by default

```
admissions <- read_csv("https://stats.idre.ucla.edu/stat/data/binary.csv")
head(admissions)
```

```
## # A tibble: 6 x 4
##   admit gre  gpa rank
##   <dbl> <dbl> <dbl> <dbl>
## 1     0  380  3.61     3
## 2     1  660  3.67     3
## 3     1  800   4     1
## 4     1  640  3.19     4
## 5     0  520  2.93     4
## 6     1  760   3     2
```

Logistic regression

```
m3<-glm(admit ~ gre + gpa + factor(rank),  
        data = admissions,  
        family = "binomial")  
tidy(m3)
```

```
## # A tibble: 6 x 5
```

##	term	estimate	std.error	statistic	p.value
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	(Intercept)	-3.99	1.14	-3.50	0.000465
## 2	gre	0.00226	0.00109	2.07	0.0385
## 3	gpa	0.804	0.332	2.42	0.0154
## 4	factor(rank)2	-0.675	0.316	-2.13	0.0328
## 5	factor(rank)3	-1.34	0.345	-3.88	0.000104
## 6	factor(rank)4	-1.55	0.418	-3.71	0.000205

Binomial count models

```
data("UCBadmit")  
head(UCBadmit)
```

##	dept	applicant.gender	admit	reject	applications
## 1	A	male	512	313	825
## 2	A	female	89	19	108
## 3	B	male	353	207	560
## 4	B	female	17	8	25
## 5	C	male	120	205	325
## 6	C	female	202	391	593

Binomial count models

```
m4<-glm(cbind(admit, reject) ~ applicant.gender +  
        factor(dept),  
        data = UCBadmit, family = "binomial")  
tidy(m4)
```

```
## # A tibble: 7 x 5
```

##	term	estimate	std.error	statistic	p.value
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	(Intercept)	0.682	0.0991	6.88	5.97e-12
## 2	applicant.gendermale	-0.0999	0.0808	-1.24	2.17e- 1
## 3	factor(dept)B	-0.0434	0.110	-0.395	6.93e- 1
## 4	factor(dept)C	-1.26	0.107	-11.8	2.41e-32
## 5	factor(dept)D	-1.29	0.106	-12.2	2.05e-34
## 6	factor(dept)E	-1.74	0.126	-13.8	2.86e-43
## 7	factor(dept)F	-3.31	0.170	-19.5	2.80e-84

```
fe<-read_csv("./fe_demo.csv")  
head(fe)
```

```
## # A tibble: 6 x 5  
##   fips state  pop deaths race.ethn  
##   <dbl> <chr> <dbl>   <dbl> <chr>  
## 1  1001 AL      483       0 latino  
## 2  1003 AL     3218       0 latino  
## 3  1005 AL      473       0 latino  
## 4  1007 AL      163       0 latino  
## 5  1009 AL     1713       0 latino  
## 6  1011 AL      359       0 latino
```

Poisson models

```
m5<-glm(deaths ~ race.ethn,  
        data = fe, family = "poisson")  
tidy(m5)
```

```
## # A tibble: 3 x 5
```

##	term	estimate	std.error	statistic	p.value
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	(Intercept)	-0.569	0.0237	-24.0	3.42e-127
## 2	race.ethnlatino	-0.441	0.0379	-11.6	2.88e- 31
## 3	race.ethnwhite	0.520	0.0299	17.4	1.23e- 67

Poisson models with offset

```
fe<-fe %>%  
  filter(pop>0)  
m6<-glm(deaths ~ race.ethn,  
        offset = log(pop),  
        data = fe, family = "poisson")  
tidy(m6)
```

```
## # A tibble: 3 x 5
```

##	term	estimate	std.error	statistic	p.value
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	(Intercept)	-8.95	0.0237	-378.	0.
## 2	race.ethnlatino	-0.732	0.0379	-19.3	3.29e-83
## 3	race.ethnwhite	-1.21	0.0299	-40.4	0.

- `MASS::glm.nb()` for negative binomial regression
- `pscl::zeroinfl()` for zero-inflated poisson regression
- `nnet::multinom()` for multinomial (categorical) regression

Frequentist multilevel models

```
library(lme4)
```

- lme4 syntax follows basic R formula syntax, but now adds variable intercept and slope terms
- basic format: `outcome ~ predictors + (1|varying slope) + (varying slope | varying intercept), data = data`

Estimating a multilevel model: country intercepts

```
library(gapminder)
m7<-lmer(lifeExp ~ year + (1|country),
        data = gapminder)
tidy(m7)
```

```
## # A tibble: 4 x 5
```

##	term	estimate	std.error	statistic	group
##	<chr>	<dbl>	<dbl>	<dbl>	<chr>
## 1	(Intercept)	-586.	10.0	-58.6	fixed
## 2	year	0.326	0.00503	64.8	fixed
## 3	sd_(Intercept).country	11.1	NA	NA	country
## 4	sd_Observation.Residual	3.58	NA	NA	Residual

Estimating a multilevel model: country intercepts, nested within continent intercepts

```
m8<-lmer(lifeExp ~ year + (1|continent/country),  
         data = gapminder)  
tidy(m8)
```

```
## # A tibble: 5 x 5  
##   term                                estimate std.error statistic group  
##   <chr>                                <dbl>     <dbl>     <dbl> <chr>  
## 1 (Intercept)                       -582.      10.9      -53.2 fixed  
## 2 year                               0.326     0.00503     64.8 fixed  
## 3 sd_(Intercept).country:continent  6.48      NA         NA    country:contine~  
## 4 sd_(Intercept).continent          9.86      NA         NA    continent  
## 5 sd_Observation.Residual           3.58      NA         NA    Residual
```

Varying slopes, country and continent intercepts

```
m9<-lmer(lifeExp ~ year + (year|continent/country),  
        data = gapminder)  
tidy(m9)
```

```
## # A tibble: 9 x 5
```

##	term	estimate	std.error	statistic	group
##	<chr>	<dbl>	<dbl>	<dbl>	<chr>
## 1	(Intercept)	-587.	10.7	-55.0	fixed
## 2	year	0.329	1.65	0.200	fixed
## 3	sd_(Intercept).country:continent	3.12	NA	NA	country:conti~
## 4	sd_year.country:continent	0.00488	NA	NA	country:conti~
## 5	cor_(Intercept).year.country:co~	-0.999	NA	NA	country:conti~
## 6	sd_(Intercept).continent	7.98	NA	NA	continent
## 7	sd_year.continent	3.68	NA	NA	continent
## 8	cor_(Intercept).year.continent	0.00249	NA	NA	continent
## 9	sd_Observation.Residual	3.55	NA	NA	Residual

GLMs and lme4, Poisson model with state intercepts

```
m10<-glmer(deaths ~ race.ethn + (1|state),  
            data = fe, family = "poisson")  
tidy(m10)
```

```
## # A tibble: 4 x 6  
##   term                estimate std.error statistic  p.value group  
##   <chr>                <dbl>    <dbl>    <dbl>    <dbl> <chr>  
## 1 (Intercept)        -0.692    0.142     -4.86  1.17e- 6 fixed  
## 2 race.ethnlatino    -0.465    0.0378   -12.3   8.44e-35 fixed  
## 3 race.ethnwhite      0.491    0.0299    16.4   1.49e-60 fixed  
## 4 sd_(Intercept).state 0.995    NA        NA      NA      state
```

Add an offset

```
m11<-glmer(deaths ~ race.ethn + offset(log(pop)) +  
            (1|state),  
            data = fe, family = "poisson")  
tidy(m11)
```

```
## # A tibble: 4 x 6
```

##	term	estimate	std.error	statistic	p.value	group
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>
## 1	(Intercept)	-8.86	0.0708	-125.	0.	fixed
## 2	race.ethnlatino	-0.997	0.0408	-24.5	3.39e-132	fixed
## 3	race.ethnwhite	-1.30	0.0314	-41.4	0.	fixed
## 4	sd_(Intercept).state	0.456	NA	NA	NA	state

Add overdispersion

We can model overdispersion in counts with a variable intercept for each observation. This effectively creates an error term for a model that typically lacks one.

```
fe<-fe %>%  
  mutate(obs_n = 1:nrow(fe))
```

Add overdispersion

```
m12<-glmer(deaths ~ race.ethn + offset(log(pop)) +  
            (1|state) + (1|obs_n),  
            data = fe, family = "poisson")  
tidy(m12)
```

```
## # A tibble: 5 x 6
```

##	term	estimate	std.error	statistic	p.value	group
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>
## 1	(Intercept)	-9.02	0.0747	-121.	0.	fixed
## 2	race.ethnlatino	-0.976	0.0572	-17.1	3.25e- 65	fixed
## 3	race.ethnwhite	-1.19	0.0417	-28.4	5.08e-178	fixed
## 4	sd_(Intercept).obs_n	0.418	NA	NA	NA	obs_n
## 5	sd_(Intercept).state	0.446	NA	NA	NA	state

Bayesian models with lme4 and glm syntax

- brms uses the core R and lme4 formula syntax, but uses rstan and HMC to fit models
- we can specify priors relatively easily
- For a full translation of Statistical Rethinking into brms, see this very detailed guide: https://bookdown.org/ajkurz/Statistical_Rethinking_recoded/


```
library(brms)
brm0<-brm(Divorce ~ Marriage,
          data = WaffleDivorce,
          prior = c(prior(normal(0,2), class = Intercept),
                    prior(normal(0,2), class = b, coef = Marriage),
                    prior(exponential(1), class = sigma))
          )
```

Results

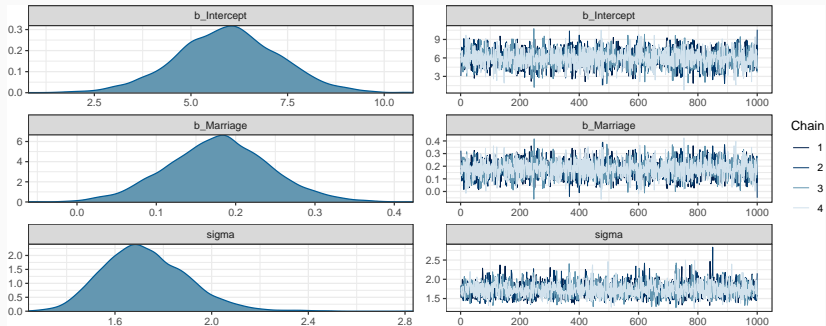
```
summary(brm0)
```

```
## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: Divorce ~ Marriage
## Data: WaffleDivorce (Number of observations: 50)
## Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##           total post-warmup samples = 4000
##
## Population-Level Effects:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept      5.95      1.33      3.28      8.57 1.00      3499      2752
## Marriage        0.18      0.06      0.05      0.31 1.00      3438      2746
##
## Family Specific Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma        1.73      0.18      1.43      2.10 1.00      3957      2965
##
## Samples were drawn using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

Traceplots and density plots

so easy!

```
plot(brm0)
```



Directly sampling the posterior

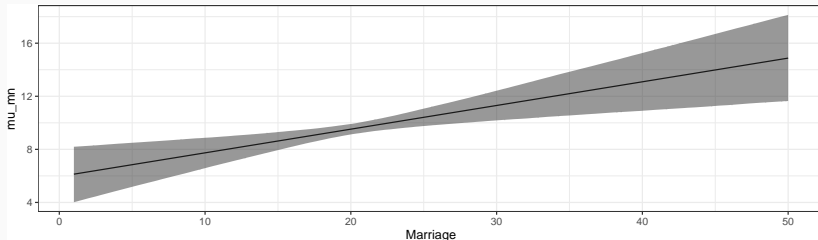
```
post<-posterior_samples(brm0)
head(post)
```

```
##      b_Intercept b_Marriage      sigma      lp__
## 1      6.077008  0.1755933  1.803016 -112.9486
## 2      5.806530  0.1884207  1.524815 -113.0423
## 3      5.985369  0.1850693  1.736283 -112.9152
## 4      4.675092  0.2417333  1.630557 -113.1536
## 5      5.894895  0.1619366  1.567004 -114.6464
## 6      5.973628  0.1716880  1.990871 -114.0798
```

Using the linear link function and new data

```
sim_dat<-data.frame(Marriage = seq(1, 50, by = 0.1))
post_mu<-posterior_linpred(brm0, newdata = sim_dat)

sim_dat<-sim_dat %>%
  mutate(mu_lwr = apply(post_mu, 2, function(x) quantile(x, 0.05)),
         mu_upr = apply(post_mu, 2, function(x) quantile(x, 0.95)),
         mu_mn = apply(post_mu, 2, mean))
ggplot(sim_dat, aes(x = Marriage, y = mu_mn)) +
  geom_line() +
  geom_ribbon(aes(ymin = mu_lwr, ymax = mu_upr), alpha = 0.5)
```



```
brm1<-brm(admit ~ gre + gpa,  
          family = bernoulli,  
          data = admissions,  
          prior = c(prior(normal(0.5, 1), class = Intercept),  
                    prior(normal(0, 1), class = b)))
```

```
summary(brm1)
```

```
## Family: bernoulli
## Links: mu = logit
## Formula: admit ~ gre + gpa
## Data: admissions (Number of observations: 400)
## Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##          total post-warmup samples = 4000
##
## Population-Level Effects:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept    -4.75      1.02   -6.80   -2.78 1.00     2468     2461
## gre           0.00      0.00    0.00    0.00 1.00     4176     3002
## gpa           0.68      0.30    0.10    1.25 1.00     2121     1811
##
## Samples were drawn using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

- brms is built to accept the syntax of lme4 - like formulas

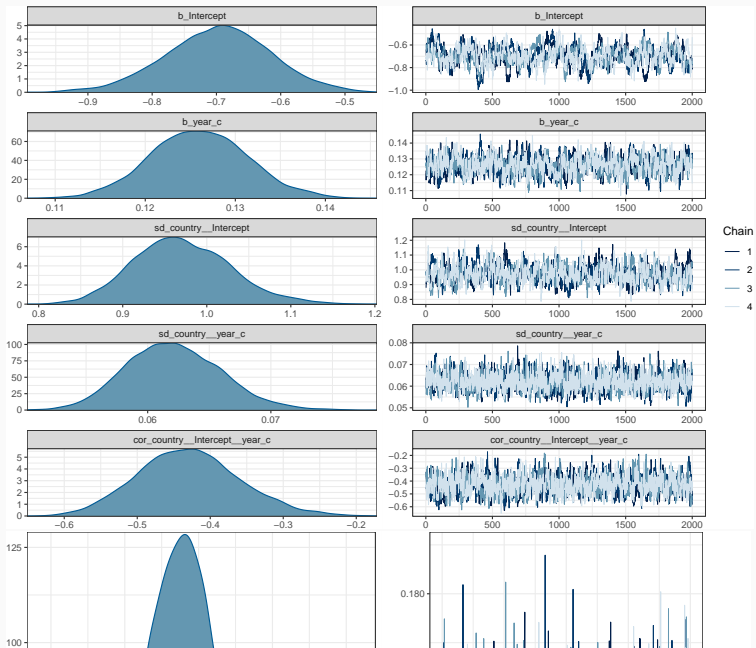
```
gapminder<-gapminder %>%  
  mutate(L_c = scale(lifeExp),  
         year_c = ((year - min(year))/5))  
brm7<-brm(L_c ~ year_c + (year_c|country),  
  data = gapminder, family = gaussian,  
  prior = c(prior(normal(0, 5), class = Intercept),  
            prior(exponential(1), class = sigma),  
            prior(normal(0,2), class = b),  
            prior(lkj(2), class = cor)),  
  iter = 4000, cores = 4)
```


Output

```
summary(brm7)
```

```
## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: L_c ~ year_c + (year_c | country)
## Data: gapminder (Number of observations: 1704)
## Samples: 4 chains, each with iter = 4000; warmup = 2000; thin = 1;
##           total post-warmup samples = 8000
##
## Group-Level Effects:
## ~country (Number of levels: 142)
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS
## sd(Intercept)      0.97      0.06      0.87      1.09 1.01      399
## sd(year_c)         0.06      0.00      0.06      0.07 1.00      907
## cor(Intercept,year_c) -0.43      0.07     -0.56     -0.28 1.01      681
##           Tail_ESS
## sd(Intercept)      1034
## sd(year_c)         1937
## cor(Intercept,year_c) 1529
##
## Population-Level Effects:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept     -0.70      0.08     -0.86     -0.54 1.01      224      421
## year_c         0.13      0.01      0.12      0.14 1.00      524     1158
##
## Family Specific Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma         0.17      0.00      0.16      0.18 1.00     8406     6097
##
```

Check convergence



- An Introduction to Bayesian Multilevel Models Using brms: A Case Study of Gender Effects on Vowel Variability in Standard Indonesian
<https://osf.io/dpzcb/>
- Fitting Linear Mixed-Effects Models Using lme4: <https://cran.r-project.org/web/packages/lme4/vignettes/lmer.pdf>
- brms: An R Package for Bayesian Multilevel Models using Stan:
https://cran.r-project.org/web/packages/brms/vignettes/brms_overview.pdf