

Measurement error and missing data

Frank Edwards

5/1/2020

What is measurement error?

- When randomly sampling a unit from a larger population, we obtain a single value
- This single value x differs by an unknown amount from the true population value, μ . This is called sampling error.
- This means that our observation is but one of many possible values that could have been observed with a random sample of the population
- Treating the observation as the **truth** is thus unwise

Divorce data again, now with measurement error!

```
library(rethinking)
```

```
data(WaffleDivorce)
```


```
d <- WaffleDivorce
```

```
head(d %>% select(Location, Divorce, Divorce.SE))
```

##	Location	Divorce	Divorce.SE
## 1	Alabama	12.7	0.79
## 2	Alaska	12.5	2.05
## 3	Arizona	10.8	0.74
## 4	Arkansas	13.5	1.22
## 5	California	8.0	0.24
## 6	Colorado	11.6	0.94

Visualizing the uncertainty


```
ggplot(d, aes(x = reorder(Location, Divorce), y = Divorce,  
  ymin = Divorce - Divorce.SE * 1.65, ymax = Divorce + Divorce.SE * 1.65)) +  
  coord_flip() +  
  geom_pointrange(size = 0.1) +  
  labs(x = "State")
```



Error is a function of population size

The American Community Survey (ACS) samples a fixed percent of the population. Places with smaller populations have fewer sampled households and larger error

```
ggplot(d, aes(x = Population, y = Divorce.SE)) +  
  geom_point()
```



What this means in practice

```
NJ<-d %>% filter(Loc=="NJ") %>% select(Loc, Divorce, Divorce.SE)
NJ
```

```
##   Loc Divorce Divorce.SE
## 1  NJ      6.1      0.46
```

Let's assume that NJ's true Divorce rate is in fact 6.1 (we got very luck in sampling!). Here's ten possible values we *could* have obtained instead of 6.1


```
rnorm(10, 6.1, 0.46)
```

```
## [1] 5.811831 6.184476 5.715611 6.833829 6.251574 5.722585 6.324217 6.439629
## [9] 6.364859 5.959521
```

What this means in practice

Let's say that the true divorce rate in NJ is 6.3, this is the density of possible observations with the one percent ACS sample.

```
samples<-rnorm(1e5, 6.3, 0.46)  
plot(density(samples))
```



Treating the truth as a parameter

For each observed value of a state's divorce rate based on a sample, D_{obs_i} , there is a true state divorce rate D_{true_i} . The observation is one draw from the true sampling distribution

$$D_{\text{obs}_i} \sim \text{Normal}(D_{\text{true}_i}, D_{\text{SE}_i})$$

We don't observe D_{true} , but can estimate a posterior for it from our model, and incorporate uncertainty in measurement across the model.

Defining a model

- Note that we'll now be estimating uncertainty in divorce rates as we estimate our model of divorce rates!

$$D_{\text{obs}_i} \sim \text{Normal}(D_{\text{true}_i}, D_{\text{SE}_i})$$

$$D_{\text{true}_i} \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_A A_i + \beta_M M_i$$

$$\alpha \sim \text{Normal}(0, 1)$$

$$\beta_A \sim \text{Normal}(0, 1)$$

$$\beta_M \sim \text{Normal}(0, 1)$$

$$\sigma \sim \text{Exponential}(1)$$


Estimating the model

```
d_slim<-list(  
  D_obs = scale(d$Divorce), A = scale(d$MedianAgeMarriage),  
  M = scale(d$Marriage), D_se = d$Divorce.SE / sd(d$Divorce),  
  N = nrow(d))  
  
m_error<-ulam(alist(  
  D_obs ~ dnorm(D_true, D_se),  
  vector[N]:D_true ~ dnorm(mu, sigma),  
  mu<-a + bA * A + bM * M,  
  a ~ dnorm(0, 0.2),  
  bA ~ dnorm(0, 0.5),  
  bM ~ dnorm(0, 0.5),  
  sigma ~ dexp(1)  
), data=d_slim, cores = 4, chains = 4)
```


Estimate a model with no measurement error adjustment for comparison

```
m_no_error<-ulam(alist(  
  D_obs ~ dnorm(mu, sigma),  
  mu<-a + bA * A + bM * M,  
  a ~ dnorm(0, 0.2),  
  bA ~ dnorm(0, 0.5),  
  bM ~ dnorm(0, 0.5),  
  sigma ~ dexp(1)  
) , data=d_slim, cores = 4, chains = 4)
```


Visualize the results




Visualize the results: error model posterior μ




Visualize the results: observed data (blue)




Visualize the results: model without error adjustment (red)



Visualize the results: magnitude of shrinkage




The posterior for each observation of D



Fitting this model with brms

```
m_error_brm<-brm(D_obs|mi(D_se) ~ A + M,  
  prior = c(  
    prior(normal(0, 0.2), class = Intercept),  
    prior(normal(0, 0.5), class = b),  
    prior(exponential(1), class = sigma)),  
  data = d_slim, save_mevars = T)
```

Visualizing the brms inferences



Measurement error on both
outcomes and predictors


Let's look at that data again

```
head(d %>%  
  select(Loc, Marriage, Marriage.SE,  
         Divorce, Divorce.SE))
```

##	Loc	Marriage	Marriage.SE	Divorce	Divorce.SE
## 1	AL	20.2	1.27	12.7	0.79
## 2	AK	26.0	2.93	12.5	2.05
## 3	AZ	20.3	0.98	10.8	0.74
## 4	AR	26.4	1.70	13.5	1.22
## 5	CA	19.1	0.39	8.0	0.24
## 6	CO	23.5	1.24	11.6	0.94

What this means in practice

If we assume that the observed marriage and divorce values are the true value, the sampling distributions look like this



Incorporating two kinds of measurement error into our model

$$D_{\text{obs}_i} \sim \text{Normal}(D_{\text{true}_i}, D_{\text{SE}_i})$$

$$D_{\text{true}_i} \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_A A_i + \beta_M M_{\text{true}_i}$$

$$M_{\text{obs}_i} \sim \text{Normal}(M_{\text{true}_i}, M_{\text{SE}_i})$$

$$M_{\text{true}_i} \sim \text{Normal}(0, 1)$$

$$\alpha \sim \text{Normal}(0, 1)$$

$$\beta_A \sim \text{Normal}(0, 1)$$

$$\beta_M \sim \text{Normal}(0, 1)$$

$$\sigma \sim \text{Exponential}(1)$$

Estimating the model with ulam()

```
d_slim$M_se<-d$Marriage.SE/sd(d$Marriage)
m_error_both<-ulam(alist(
  D_obs ~ dnorm(D_true, D_se),
  vector[N]:D_true ~ dnorm(mu, sigma),
  mu<-a + bA * A + bM * M,
  M ~ dnorm(M_true, M_se),
  vector[N]:M_true ~ dnorm(0,1),
  a ~ dnorm(0, 0.2),
  bA ~ dnorm(0, 0.5),
  bM ~ dnorm(0, 0.5),
  sigma ~ dexp(1)),
  data=d_slim, cores = 4, chains = 4)
```


Estimating the model with brm()

```
m_error_brm_both<-brm(D_obs|mi(D_se) ~ A +  
                      me(M, M_se),  
                      prior = c(  
                        prior(normal(0, 0.2), class = Intercept),  
                        prior(normal(0, 0.5), class = b),  
                        prior(exponential(1), class = sigma)),  
                      data = d_slim, save_mevars = T)
```

- Data are often measured with error
- Sometimes we know this error - that's great!
- Incorporate the measurement error into your model
- Bayesian models are generative: we can estimate a posterior for variables measured with error at the same time as we estimate other parameters
- Correlated errors across multiple measures can induce bias (see DAGs on page 498)

Missing data

- Sometimes variables are missing a value

- Sometimes variables are missing a value
- R calls these **NA**

- Sometimes variables are missing a value
- R calls these **NA**
- There can be many reasons that a value isn't reported

- Sometimes variables are missing a value
- R calls these **NA**
- There can be many reasons that a value isn't reported
- Think hard about why it may have happened

- Sometimes variables are missing a value
- R calls these **NA**
- There can be many reasons that a value isn't reported
- Think hard about why it may have happened
- Software typically defaults to *listwise deletion*, or *complete case analysis*

- Sometimes variables are missing a value
- R calls these **NA**
- There can be many reasons that a value isn't reported
- Think hard about why it may have happened
- Software typically defaults to *listwise deletion*, or *complete case analysis*
- At best, this discards perfectly good information in the other variables in that row. At worst, it leads to biased inference.

- Avoid discarding data whenever possible

- Avoid discarding data whenever possible
- Use common sense to replace missings when possible (e.g. in panel data when we have a respondent's age in one wave but not another)

- Avoid discarding data whenever possible
- Use common sense to replace missings when possible (e.g. in panel data when we have a respondent's age in one wave but not another)
- Avoid imputing data with a single value when we are uncertain (never replace a missing with a mean or interpolation). This overstates your certainty, and artificially narrows posterior intervals

- Avoid discarding data whenever possible
- Use common sense to replace missings when possible (e.g. in panel data when we have a respondent's age in one wave but not another)
- Avoid imputing data with a single value when we are uncertain (never replace a missing with a mean or interpolation). This overstates your certainty, and artificially narrows posterior intervals
- Use imputation models to populate your uncertainty into the analysis.

The ludicrous names that get used in the literature to describe various mechanisms for missing data

- Missing completely at random: each observation has equal probability of being missing.
- Missing at random: each observation's probability of being missing is conditional on some set of measured variables
- Missing not at random: each observation's probability of being missing is conditional on some set of unmeasured variables

This dog ate my homework

He's still a good boy




Why do dogs make homework go missing?

Let's assume that H is the grade a homework would have received if graded, S is how much a student studied, and D is whether a dog ate the homework. We observe H_{obs} , a vector of grades where some are missing.

Dogs eat homework completely at random

Nothing affects a dog's decision to eat homework, they just strike at random!



Simulate random homework eating

Dogs eat homework 20 percent of the time, and strike at random

```
N<-100
S<-rnorm(N)
H<-rbinom(N, size = 10, prob = inv_logit(S))
D<-rbinom(N, 1, p = 0.2)
H_obs<-H
H_obs[D==1]<-NA
summary(H_obs)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	0.000	3.000	4.000	4.692	6.000	10.000	22

How does the missingness affect inference? List-wise deletion

```
library(broom)
d<-data.frame(H_obs = H_obs, S = S) %>% filter(!is.na(H_obs))
mcar_0<-brm(H_obs ~ S, data = d)
```

```
tidy(mcar_0)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	5.010457	0.1728653	4.724120	5.295409
## 2	b_S	2.367944	0.2097332	2.020992	2.718633
## 3	sigma	1.480677	0.1239770	1.290049	1.696584
## 4	lp__	-145.912882	1.2764912	-148.397932	-144.559153

How does the missingness affect inference? Complete data

```
d<-data.frame(H = H, S = S)
mcar_1<-brm(H ~ S, data = d)
```

```
tidy(mcar_1)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	5.056265	0.1428282	4.817824	5.287313
## 2	b_S	2.308409	0.1651645	2.045640	2.589497
## 3	sigma	1.411133	0.1045346	1.251116	1.593560
## 4	lp__	-181.274366	1.2884789	-183.813197	-179.913627

Dogs eat homework when students study too much

Why don't you want to go play?



Simulate bad dogs

```
P<-ifelse(S>0, 0.5, 0)
D<-rbinom(N, 1, prob = P)
H_obs<-H
H_obs[D==1]<-NA
summary(H_obs)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	0.000	3.000	4.000	4.273	6.000	10.000	23

How does the missingness affect inference? List-wise deletion

```
d<-data.frame(H_obs = H_obs, S = S) %>% filter(!is.na(H_obs))
mar_0<-brm(H_obs ~ S, data = d)
```

```
tidy(mar_0)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	5.123431	0.1615340	4.853224	5.389861
## 2	b_S	2.393996	0.1817246	2.094337	2.696579
## 3	sigma	1.319999	0.1105789	1.147163	1.511831
## 4	lp__	-135.513357	1.2176541	-137.888184	-134.165899

How does the missingness affect inference? Complete data


```
d<-data.frame(H = H, S = S)
mar_1<-brm(H ~ S, data = d)
```

```
tidy(mar_1)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	5.057506	0.1429910	4.825733	5.290436
## 2	b_S	2.309885	0.1660931	2.043301	2.585037
## 3	sigma	1.412079	0.1015118	1.256419	1.584946
## 4	lp__	-181.251686	1.2327416	-183.704430	-179.905489

Dogs hate noise, and so does homework

Noisy homes X make bad homework and dogs



Simulate noisy houses

```
N<-100
S<-rnorm(N)
X<-rnorm(N)
H<-rbinom(N, size = 10, prob = inv_logit(2 + S - 2*X))
D<-ifelse(X>1, 1, 0)
H_obs<-H
H_obs[D==1]<-NA
summary(H_obs)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	2.000	8.000	9.000	8.247	10.000	10.000	19

How does the missingness affect inference? List-wise deletion

```
d<-data.frame(H_obs = H_obs, S = S, X = X) %>% filter(!is.na(H_obs))  
mnar_0<-brm(H_obs ~ S, data = d)
```

```
tidy(mnar_0)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	8.377476	0.2212538	8.0167669	8.737188
## 2	b_S	1.153430	0.2229636	0.7812635	1.530252
## 3	sigma	1.957408	0.1588300	1.7158168	2.239891
## 4	lp__	-173.878775	1.2765405	-176.3685728	-172.527316

How does the missingness affect inference? Complete data

```
d<-data.frame(H = H, S = S)
mnar_1<-brm(H ~ S, data = d)
```

```
tidy(mnar_1)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	7.360858	0.2998019	6.877486	7.858758
## 2	b_S	1.506460	0.3009301	1.020529	2.007309
## 3	sigma	3.011228	0.2237737	2.663287	3.404937
## 4	lp__	-256.092820	1.2324443	-258.498001	-254.743549

How does the missingness affect inference? Complete data with unobserved cause


```
d<-data.frame(H = H, S = S, X = X)
mnar_2<-brm(H ~ S + X, data = d)
```

```
tidy(mnar_2)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	7.457159	0.1740462	7.1747142	7.743018
## 2	b_S	1.117298	0.1772999	0.8296919	1.404930
## 3	b_X	-2.304900	0.1615164	-2.5704538	-2.039863
## 4	sigma	1.695863	0.1210468	1.5112663	1.911102
## 5	lp__	-199.693354	1.3923707	-202.4171385	-198.019594

Very good dogs

Dogs only eat bad homework



Simulate good dogs

```
N<-100
S<-rnorm(N)
H<-rbinom(N, size = 10, prob = inv_logit(S))
D<-ifelse(H<5, 1, 0)
H_obs<-H
H_obs[D==1]<-NA
summary(H_obs)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	5.000	5.000	6.000	6.738	8.000	10.000	39

How does the missingness affect inference? List-wise deletion

```
d<-data.frame(H_obs = H_obs, S = S) %>% filter(!is.na(H_obs))  
mnar_0<-brm(H_obs ~ S, data = d)
```

```
tidy(mnar_0)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	5.754663	0.2103699	5.4076133	6.105897
## 2	b_S	1.515797	0.2304014	1.1362693	1.899124
## 3	sigma	1.145530	0.1057819	0.9897332	1.332844
## 4	lp__	-99.865863	1.1903074	-102.2037173	-98.517445

How does the missingness affect inference? Complete data

```
d<-data.frame(H = H, S = S)
mnar_1<-brm(H ~ S, data = d)
```

```
tidy(mnar_1)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	4.903829	0.1389065	4.674313	5.128918
## 2	b_S	2.278978	0.1520660	2.023544	2.531579
## 3	sigma	1.405604	0.1040351	1.245806	1.591546
## 4	lp__	-180.754545	1.2672585	-183.300258	-179.401340

- Much like with measurement error, we can treat missing values as parameters to be estimated from our model.
- HMC samples continuous missing measures well
- But it can't handle categorical measures as easily (see 15.3 for a method)
- We're going to use `brm()`, but see 15.2.2 for examples using `ulam()`

Bayesian imputation with brm()

```
library(mice)
data(nhanes)
summary(nhanes)
```

##	age	bmi	hyp	chl
##	Min. :1.00	Min. :20.40	Min. :1.000	Min. :113.0
##	1st Qu.:1.00	1st Qu.:22.65	1st Qu.:1.000	1st Qu.:185.0
##	Median :2.00	Median :26.75	Median :1.000	Median :187.0
##	Mean :1.76	Mean :26.56	Mean :1.235	Mean :191.4
##	3rd Qu.:2.00	3rd Qu.:28.93	3rd Qu.:1.000	3rd Qu.:212.0
##	Max. :3.00	Max. :35.30	Max. :2.000	Max. :284.0
##		NA's :9	NA's :8	NA's :10

Let's build a model to predict BMI as a function of age and cholesterol

```
m0<-brm(bmi ~ age + chl,  
        data = nhanes, cores = 4)
```

```
tidy(m0)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	19.41425783	4.08673940	12.77672986	26.0945251
## 2	b_age	-5.34252093	1.68150558	-8.01861983	-2.6468062
## 3	b_chl	0.08389625	0.02576278	0.04317525	0.1262215
## 4	sigma	3.45781357	0.89396538	2.35098064	5.0549389
## 5	lp__	-38.53672426	1.63118244	-41.71180431	-36.6707266

This model dropped the missing values in bmi and chl

- We can impute values by sampling them with HMC, giving them a likelihood and priors
- We'll obtain posteriors for each missing value
- We specify these likelihoods and priors explicitly

```
m1 <- brm(bmi|mi() ~ age + chl,  
          data = nhanes)
```

But we've got missingness on chl too

We need to specify a model with more than one likelihood, as we did in the measurement error model. We set (missing and non-missing) values in bmi to be a function of age and cholesterol, and missing cholesterol values to be a function of age.

```
formula_2 <-  
  bf(bmi | mi() ~ age + mi(chl)) + # BMI model  
  bf(chl | mi() ~ age) # chl model  
m2 <- brm(formula_2,  
           data = nhanes)
```


The output

```
summary(m2)
```

```
## Family: MV(gaussian, gaussian)
## Links: mu = identity; sigma = identity
##      mu = identity; sigma = identity
## Formula: bmi | mi() ~ age + mi(chl)
##      chl | mi() ~ age
## Data: nhanes (Number of observations: 25)
## Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##      total post-warmup samples = 4000
##
## Population-Level Effects:
##      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## bmi_Intercept    21.94     3.82   14.65   29.84 1.00    1760    2101
## chl_Intercept    138.72    24.74   89.72  186.90 1.00    2654    2774
## bmi_age          -4.29     1.48   -7.17   -1.34 1.00    1274    2240
## chl_age           30.84    12.90    5.73   56.47 1.00    2707    2650
## bmi_michl         0.06     0.02    0.01    0.11 1.00    1070    1526
##
## Family Specific Parameters:
##      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma_bmi       3.31     0.77    2.15    5.08 1.00    1558    1909
## sigma_chl       40.18     7.65   28.27   57.83 1.00    2095    2085
##
## Samples were drawn using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

Visualizing the imputation

```
post<-posterior_samples(m2)  
plot(density(post$Ymi_bmi[10]))
```



Pros - We can allow for very complex structures in our models - We can neatly specify variable specific models for missing data - Sampling missing values using HMC as part of the model is consistent with Bayesian modeling principals

Cons - Gets technically complex quickly - Doesn't have easy solution for categorical data - Computationally intensive for big data

Multiple imputation by chained equations (MICE)

MICE is a common pseudo-Bayesian approach to missing data. It produces k predictions for each unobserved variable based on a fully conditional regression models (where each variable is a function of all others).

- We'll allow each variable in the model to be a function of all others (we can relax this)
- Then make conditional predictions for each missing value
- This results in k complete datasets
- We apply our analysis over each dataset, and pool results

For a detailed description of the method and software, See Van Buuren and Groothuis-Oudshoorn, 2011:


<http://www.jstatsoft.org/v45/i03/paper>

It's honestly too easy

```
library(mice)
imp_nhanes<-mice(nhanes, m = 10, maxit = 20)
summary(imp_nhanes)
```


Visualizing convergence

```
plot(imp_nhanes)
```



Visualizing the imputations

```
stripplot(imp_nhanes, chl~.imp, pch=20, cex=2)
```



Changing predictors

If we've got a collinear variable or a factor with many categories, we may want to exclude it from the model. Columns indicate predictors, row indicates outcomes.

```
pred<-imp_nhanes$predictorMatrix
```

Turn off hyp as a predictor of chl

```
pred["chl", "hyp"]<-0  
pred
```

```
##      age bmi hyp chl  
## age    0   1   1   1  
## bmi    1   0   1   1  
## hyp    1   1   0   1  
## chl    1   1   0   0
```

Use the new predictor matrix

```
imp_nhanes2<-mice(nhanes, m = 10, maxit = 20,  
                  predictorMatrix = pred)  
summary(imp_nhanes)
```

Change imputation methods

- The default in mice for continuous measures is partial mean matching
- The algorithm constructs a regression model, samples new coefficients from the parameter distribution, then makes a prediction. It then randomly selects one proximate observation from the observed data
- This makes the imputation data appear similar to the observed
- But sometimes we may not want to do this, or may have categorical data

Imputation methods

For a full list of methods, see

<https://stefvanbuuren.name/fimd/sec-modelform.html>

```
meth<-imp_nhanes$method
```

Let's change chl to a linear model


```
meth[4]<-"norm"
```

And re-run the model

```
imp_nhanes3<-mice(nhanes, m = 10, maxit = 20,  
                  method = meth)  
summary(imp_nhanes3)
```

Visualize imputations

```
densityplot(imp_nhanes3)
```



Using the imputed data

After imputation, we conduct our analysis over each imputed dataset.

Typically, we'll need to separately fit each model.

For frequentist methods, either use a loop or the `with()` function

```
fit <- with(imp_nhanes3, lm(bmi ~ chl))
```

Then pool your results according to Rubin's rules for combination. This averages beta parameters, and adjusts standard errors for cross-imputation variance

The results

```
pool(fit)
```

```
## Class: mipo    m = 10
##      term  m    estimate      ubar      b      t dfcom
## 1 (Intercept) 10 21.31695134 1.113277e+01 2.129882e+00 1.347564e+01 23
## 2      chl 10  0.02587948 2.925939e-04 5.994597e-05 3.585345e-04 23
##      df      riv  lambda      fmi
## 1 16.56385 0.2104482 0.1738597 0.2583155
## 2 16.26681 0.2253655 0.1839170 0.2686309
```

In a Bayesian context, we don't have to worry about formulas for combination. We fit the model to each dataset, then pool the posterior samples from each for inference.

brms provides `brm_multiple()` to do this

```
brm_mi <- brm_multiple(bmi ~ age + chl,  
  data = imp_nhanes3, cores = 4)
```

```
summary(brm_mi)
```

```
## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: bmi ~ age + chl
## Data: imp_nhanes3 (Number of observations: 25)
## Samples: 40 chains, each with iter = 2000; warmup = 1000; thin = 1;
##          total post-warmup samples = 40000
##
## Population-Level Effects:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept      20.85      3.18   14.26   26.84 1.18      150      634
## age             -3.61      1.01   -5.59   -1.60 1.10      235     1224
## chl              0.06      0.02    0.02    0.10 1.19      143      685
##
## Family Specific Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma          3.08      0.54    2.21    4.33 1.09      261      713
##
## Samples were drawn using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

The posterior samples

```
post<-posterior_samples(brm_mi)  
nrow(post)
```

```
## [1] 40000
```

- Measurement error is an important part of the data generating process
- Ignoring it presumes we have perfect measurement, can bias inference, and understates uncertainty
- Missing data is also important. List-wise deletion at best discards information (and overstates certainty), at worst biases inference
- If data is missing completely at random, or is conditional on an observed variable, we can impute the missings and recover valid inference
- If missingness is correlated with the outcome, we've got a problem

- Further reading on mice:
<http://www.gerkovink.com/miceVignettes/>
- Further reading on Bayesian imputation and brms:
https://cran.r-project.org/web/packages/brms/vignettes/brms_missings.html

Thank you!

Thanks for a wonderful semester and for being experimental subjects as I try teaching Bayes!

Stay safe, and please be in touch: `frank.edwards@rutgers.edu`