

# Measurement error and missing data

---

Frank Edwards

5/1/2020

## What is measurement error?

- When randomly sampling a unit from a larger population, we obtain a single value
- This single value  $x$  differs by an unknown amount from the true population value,  $\mu$ . This is called sampling error.
- This means that our observation is but one of many possible values that could have been observed with a random sample of the population
- Treating the observation as the **truth** is thus unwise

## Divorce data again, now with measurement error!

```
library(rethinking)
```

```
data(WaffleDivorce)
```


```
d <- WaffleDivorce
```

```
head(d %>% select(Location, Divorce, Divorce.SE))
```

##	Location	Divorce	Divorce.SE
## 1	Alabama	12.7	0.79
## 2	Alaska	12.5	2.05
## 3	Arizona	10.8	0.74
## 4	Arkansas	13.5	1.22
## 5	California	8.0	0.24
## 6	Colorado	11.6	0.94

# Visualizing the uncertainty


```
ggplot(d, aes(x = reorder(Location, Divorce), y = Divorce,  
  ymin = Divorce - Divorce.SE * 1.65, ymax = Divorce + Divorce.SE * 1.65)) +  
  coord_flip() +  
  geom_pointrange(size = 0.1) +  
  labs(x = "State")
```



# Error is a function of population size

The American Community Survey (ACS) samples a fixed percent of the population. Places with smaller populations have fewer sampled households and larger error

```
ggplot(d, aes(x = Population, y = Divorce.SE)) +  
  geom_point()
```



## What this means in practice

```
NJ<-d %>% filter(Loc=="NJ") %>% select(Loc, Divorce, Divorce.SE)
NJ
```

```
##   Loc Divorce Divorce.SE
## 1  NJ      6.1      0.46
```

Let's assume that NJ's true Divorce rate is in fact 6.1 (we got very luck in sampling!). Here's ten possible values we *could* have obtained instead of 6.1


```
rnorm(10, 6.1, 0.46)
```

```
## [1] 5.811831 6.184476 5.715611 6.833829 6.251574 5.722585 6.324217 6.439629
## [9] 6.364859 5.959521
```

## What this means in practice

Let's say that the true divorce rate in NJ is 6.3, this is the density of possible observations with the one percent ACS sample.

```
samples<-rnorm(1e5, 6.3, 0.46)  
plot(density(samples))
```



## Treating the truth as a parameter

For each observed value of a state's divorce rate based on a sample,  $D_{\text{obs}_i}$ , there is a true state divorce rate  $D_{\text{true}_i}$ . The observation is one draw from the true sampling distribution

$$D_{\text{obs}_i} \sim \text{Normal}(D_{\text{true}_i}, D_{\text{SE}_i})$$

We don't observe  $D_{\text{true}}$ , but can estimate a posterior for it from our model, and incorporate uncertainty in measurement across the model.



## Defining a model

- Note that we'll now be estimating uncertainty in divorce rates as we estimate our model of divorce rates!

$$D_{\text{obs}_i} \sim \text{Normal}(D_{\text{true}_i}, D_{\text{SE}_i})$$

$$D_{\text{true}_i} \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_A A_i + \beta_M M_i$$

$$\alpha \sim \text{Normal}(0, 1)$$

$$\beta_A \sim \text{Normal}(0, 1)$$

$$\beta_M \sim \text{Normal}(0, 1)$$

$$\sigma \sim \text{Exponential}(1)$$


# Estimating the model

```
d_slim<-list(  
  D_obs = scale(d$Divorce), A = scale(d$MedianAgeMarriage),  
  M = scale(d$Marriage), D_se = d$Divorce.SE / sd(d$Divorce),  
  N = nrow(d))  
  
m_error<-ulam(alist(  
  D_obs ~ dnorm(D_true, D_se),  
  vector[N]:D_true ~ dnorm(mu, sigma),  
  mu<-a + bA * A + bM * M,  
  a ~ dnorm(0, 0.2),  
  bA ~ dnorm(0, 0.5),  
  bM ~ dnorm(0, 0.5),  
  sigma ~ dexp(1)  
, data=d_slim, cores = 4, chains = 4)
```


## Estimate a model with no measurement error adjustment for comparison

```
m_no_error<-ulam(alist(  
  D_obs ~ dnorm(mu, sigma),  
  mu<-a + bA * A + bM * M,  
  a ~ dnorm(0, 0.2),  
  bA ~ dnorm(0, 0.5),  
  bM ~ dnorm(0, 0.5),  
  sigma ~ dexp(1)  
) , data=d_slim, cores = 4, chains = 4)
```


## Visualize the results




## Visualize the results: error model posterior mu




## Visualize the results: observed data (blue)



## Visualize the results: model without error adjustment (red)




## Visualize the results: magnitude of shrinkage






# The posterior for each observation of D



## Fitting this model with brms

```
m_error_brm<-brm(D_obs|mi(D_se) ~ A + M,  
  prior = c(  
    prior(normal(0, 0.2), class = Intercept),  
    prior(normal(0, 0.5), class = b),  
    prior(exponential(1), class = sigma)),  
  data = d_slim, save_mevars = T)
```

# Visualizing the brms inferences



Measurement error on both  
outcomes and predictors

---


## Let's look at that data again

```
head(d %>%  
  select(Loc, Marriage, Marriage.SE,  
         Divorce, Divorce.SE))
```

##	Loc	Marriage	Marriage.SE	Divorce	Divorce.SE
## 1	AL	20.2	1.27	12.7	0.79
## 2	AK	26.0	2.93	12.5	2.05
## 3	AZ	20.3	0.98	10.8	0.74
## 4	AR	26.4	1.70	13.5	1.22
## 5	CA	19.1	0.39	8.0	0.24
## 6	CO	23.5	1.24	11.6	0.94

## What this means in practice

If we assume that the observed marriage and divorce values are the true value, the sampling distributions look like this



## Incorporating two kinds of measurement error into our model

$$D_{\text{obs}_i} \sim \text{Normal}(D_{\text{true}_i}, D_{\text{SE}_i})$$

$$D_{\text{true}_i} \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha + \beta_A A_i + \beta_M M_{\text{true}_i}$$

$$M_{\text{obs}_i} \sim \text{Normal}(M_{\text{true}_i}, M_{\text{SE}_i})$$

$$M_{\text{true}_i} \sim \text{Normal}(0, 1)$$

$$\alpha \sim \text{Normal}(0, 1)$$

$$\beta_A \sim \text{Normal}(0, 1)$$

$$\beta_M \sim \text{Normal}(0, 1)$$

$$\sigma \sim \text{Exponential}(1)$$

## Estimating the model with ulam()

```
d_slim$M_se<-d$Marriage.SE/sd(d$Marriage)
m_error_both<-ulam(alist(
  D_obs ~ dnorm(D_true, D_se),
  vector[N]:D_true ~ dnorm(mu, sigma),
  mu<-a + bA * A + bM * M,
  M ~ dnorm(M_true, M_se),
  vector[N]:M_true ~ dnorm(0,1),
  a ~ dnorm(0, 0.2),
  bA ~ dnorm(0, 0.5),
  bM ~ dnorm(0, 0.5),
  sigma ~ dexp(1)),
  data=d_slim, cores = 4, chains = 4)
```



## Estimating the model with brm()

```
m_error_brm_both<-brm(D_obs|mi(D_se) ~ A +  
                      me(M, M_se),  
                      prior = c(  
                        prior(normal(0, 0.2), class = Intercept),  
                        prior(normal(0, 0.5), class = b),  
                        prior(exponential(1), class = sigma)),  
                      data = d_slim, save_mevars = T)
```

- Data are often measured with error
- Sometimes we know this error - that's great!
- Incorporate the measurement error into your model
- Bayesian models are generative: we can estimate a posterior for variables measured with error at the same time as we estimate other parameters
- Correlated errors across multiple measures can induce bias (see DAGs on page 498)

## Missing data

---

- Sometimes variables are missing a value

- Sometimes variables are missing a value
- R calls these **NA**

- Sometimes variables are missing a value
- R calls these **NA**
- There can be many reasons that a value isn't reported

- Sometimes variables are missing a value
- R calls these **NA**
- There can be many reasons that a value isn't reported
- Think hard about why it may have happened

- Sometimes variables are missing a value
- R calls these **NA**
- There can be many reasons that a value isn't reported
- Think hard about why it may have happened
- Software typically defaults to *listwise deletion*, or *complete case analysis*



- Sometimes variables are missing a value
- R calls these **NA**
- There can be many reasons that a value isn't reported
- Think hard about why it may have happened
- Software typically defaults to *listwise deletion*, or *complete case analysis*
- At best, this discards perfectly good information in the other variables in that row. At worst, it leads to biased inference.

- Avoid discarding data whenever possible

- Avoid discarding data whenever possible
- Use common sense to replace missings when possible (e.g. in panel data when we have a respondent's age in one wave but not another)

- Avoid discarding data whenever possible
- Use common sense to replace missings when possible (e.g. in panel data when we have a respondent's age in one wave but not another)
- Avoid imputing data with a single value when we are uncertain (never replace a missing with a mean or interpolation). This overstates your certainty, and artificially narrows posterior intervals

- Avoid discarding data whenever possible
- Use common sense to replace missings when possible (e.g. in panel data when we have a respondent's age in one wave but not another)
- Avoid imputing data with a single value when we are uncertain (never replace a missing with a mean or interpolation). This overstates your certainty, and artificially narrows posterior intervals
- Use imputation models to populate your uncertainty into the analysis.

The ludicrous names that get used in the literature to describe various mechanisms for missing data

- Missing completely at random: each observation has equal probability of being missing.
- Missing at random: each observation's probability of being missing is conditional on some set of measured variables
- Missing not at random: each observation's probability of being missing is conditional on some set of unmeasured variables

# This dog ate my homework

He's still a good boy




## Why do dogs make homework go missing?

Let's assume that  $H$  is the grade a homework would have received if graded,  $S$  is how much a student studied, and  $D$  is whether a dog ate the homework. We observe  $H_{obs}$ , a vector of grades where some are missing.



## Dogs eat homework completely at random

Nothing affects a dog's decision to eat homework, they just strike at random!



# Simulate random homework eating

Dogs eat homework 20 percent of the time, and strike at random

```
N<-100
S<-rnorm(N)
H<-rbinom(N, size = 10, prob = inv_logit(S))
D<-rbinom(N, 1, p = 0.2)
H_obs<-H
H_obs[D==1]<-NA
summary(H_obs)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	0.000	3.000	4.000	4.692	6.000	10.000	22

## How does the missingness affect inference? List-wise deletion

```
library(broom)
d<-data.frame(H_obs = H_obs, S = S) %>% filter(!is.na(H_obs))
mcar_0<-brm(H_obs ~ S, data = d)
```

```
tidy(mcar_0)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	5.010457	0.1728653	4.724120	5.295409
## 2	b_S	2.367944	0.2097332	2.020992	2.718633
## 3	sigma	1.480677	0.1239770	1.290049	1.696584
## 4	lp__	-145.912882	1.2764912	-148.397932	-144.559153

## How does the missingness affect inference? Complete data

```
d<-data.frame(H = H, S = S)
mcar_1<-brm(H ~ S, data = d)
```

```
tidy(mcar_1)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	5.056265	0.1428282	4.817824	5.287313
## 2	b_S	2.308409	0.1651645	2.045640	2.589497
## 3	sigma	1.411133	0.1045346	1.251116	1.593560
## 4	lp__	-181.274366	1.2884789	-183.813197	-179.913627

## Dogs eat homework when students study too much

Why don't you want to go play?



## Simulate bad dogs

```
P<-ifelse(S>0, 0.5, 0)
D<-rbinom(N, 1, p = P)
H_obs<-H
H_obs[D==1]<-NA
summary(H_obs)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	0.000	3.000	4.000	4.273	6.000	10.000	23

## How does the missingness affect inference? List-wise deletion

```
d<-data.frame(H_obs = H_obs, S = S) %>% filter(!is.na(H_obs))  
mar_0<-brm(H_obs ~ S, data = d)
```

```
tidy(mar_0)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	5.119811	0.1648090	4.850172	5.385539
## 2	b_S	2.391552	0.1840305	2.093521	2.694043
## 3	sigma	1.321968	0.1115854	1.151886	1.514636
## 4	lp__	-135.532848	1.2370945	-137.988813	-134.175817

## How does the missingness affect inference? Complete data

```
d<-data.frame(H = H, S = S)
mar_1<-brm(H ~ S, data = d)
```


```
tidy(mar_1)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	5.057506	0.1429910	4.825733	5.290436
## 2	b_S	2.309885	0.1660931	2.043301	2.585037
## 3	sigma	1.412079	0.1015118	1.256419	1.584946
## 4	lp__	-181.251686	1.2327416	-183.704430	-179.905489



## Dogs hate noise, and so does homework

Noisy homes  $X$  make bad homework and dogs



## Simulate noisy houses

```
N<-100
S<-rnorm(N)
X<-rnorm(N)
H<-rbinom(N, size = 10, prob = inv_logit(2 + S - 2*X))
D<-ifelse(X>1, 1, 0)
H_obs<-H
H_obs[D==1]<-NA
summary(H_obs)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	2.000	8.000	9.000	8.247	10.000	10.000	19

## How does the missingness affect inference? List-wise deletion

```
d<-data.frame(H_obs = H_obs, S = S, X = X) %>% filter(!is.na(H_obs))  
mnar_0<-brm(H_obs ~ S, data = d)
```

```
tidy(mnar_0)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	8.377476	0.2212538	8.0167669	8.737188
## 2	b_S	1.153430	0.2229636	0.7812635	1.530252
## 3	sigma	1.957408	0.1588300	1.7158168	2.239891
## 4	lp__	-173.878775	1.2765405	-176.3685728	-172.527316

## How does the missingness affect inference? Complete data

```
d<-data.frame(H = H, S = S)
mnar_1<-brm(H ~ S, data = d)
```

```
tidy(mnar_1)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	7.360858	0.2998019	6.877486	7.858758
## 2	b_S	1.506460	0.3009301	1.020529	2.007309
## 3	sigma	3.011228	0.2237737	2.663287	3.404937
## 4	lp__	-256.092820	1.2324443	-258.498001	-254.743549

# How does the missingness affect inference? Complete data with unobserved cause


```
d<-data.frame(H = H, S = S, X = X)
mnar_2<-brm(H ~ S + X, data = d)
```

```
tidy(mnar_2)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	7.457159	0.1740462	7.1747142	7.743018
## 2	b_S	1.117298	0.1772999	0.8296919	1.404930
## 3	b_X	-2.304900	0.1615164	-2.5704538	-2.039863
## 4	sigma	1.695863	0.1210468	1.5112663	1.911102
## 5	lp__	-199.693354	1.3923707	-202.4171385	-198.019594

## Very good dogs

Dogs only eat bad homework



## Simulate good dogs

```
N<-100
S<-rnorm(N)
H<-rbinom(N, size = 10, prob = inv_logit(S))
D<-ifelse(H<5, 1, 0)
H_obs<-H
H_obs[D==1]<-NA
summary(H_obs)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
##	5.000	5.000	6.000	6.738	8.000	10.000	39

## How does the missingness affect inference? List-wise deletion

```
d<-data.frame(H_obs = H_obs, S = S) %>% filter(!is.na(H_obs))  
mnar_0<-brm(H_obs ~ S, data = d)
```

```
tidy(mnar_0)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	5.754663	0.2103699	5.4076133	6.105897
## 2	b_S	1.515797	0.2304014	1.1362693	1.899124
## 3	sigma	1.145530	0.1057819	0.9897332	1.332844
## 4	lp__	-99.865863	1.1903074	-102.2037173	-98.517445



## How does the missingness affect inference? Complete data

```
d<-data.frame(H = H, S = S)
mnar_1<-brm(H ~ S, data = d)
```

```
tidy(mnar_1)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	4.903829	0.1389065	4.674313	5.128918
## 2	b_S	2.278978	0.1520660	2.023544	2.531579
## 3	sigma	1.405604	0.1040351	1.245806	1.591546
## 4	lp__	-180.754545	1.2672585	-183.300258	-179.401340

- Much like with measurement error, we can treat missing values as parameters to be estimated from our model.
- HMC samples continuous missing measures well
- But it can't handle categorical measures as easily (see 15.3 for a method)
- We're going to use `brm()`, but see 15.2.2 for examples using `ulam()`

## Bayesian imputation with brm()

```
library(mice)
data(nhanes)
summary(nhanes)
```

##	age	bmi	hyp	chl
##	Min. :1.00	Min. :20.40	Min. :1.000	Min. :113.0
##	1st Qu.:1.00	1st Qu.:22.65	1st Qu.:1.000	1st Qu.:185.0
##	Median :2.00	Median :26.75	Median :1.000	Median :187.0
##	Mean :1.76	Mean :26.56	Mean :1.235	Mean :191.4
##	3rd Qu.:2.00	3rd Qu.:28.93	3rd Qu.:1.000	3rd Qu.:212.0
##	Max. :3.00	Max. :35.30	Max. :2.000	Max. :284.0
##		NA's :9	NA's :8	NA's :10

## Let's build a model to predict BMI as a function of age and cholesterol

```
m0<-brm(bmi ~ age + chl,  
        data = nhanes, cores = 4)
```

```
tidy(m0)
```

##	term	estimate	std.error	lower	upper
## 1	b_Intercept	19.37909555	4.30609437	12.41820478	26.1394491
## 2	b_age	-5.34243116	1.81800293	-8.31855456	-2.5025320
## 3	b_chl	0.08424802	0.02803889	0.04018392	0.1298188
## 4	sigma	3.53870045	0.99612633	2.32355308	5.3745995
## 5	lp__	-38.78828222	1.83988763	-42.26928157	-36.6939884

## This model dropped the missing values in bmi and chl

- We can impute values by sampling them with HMC, giving them a likelihood and priors
- We'll obtain posteriors for each missing value
- We specify these likelihoods and priors explicitly

```
m1 <- brm(bmi|mi() ~ age + chl,  
          data = nhanes)
```

## But we've got missingness on chl too

We need to specify a model with more than one likelihood, as we did in the measurement error model. We set (missing and non-missing) values in bmi to be a function of age and cholesterol, and missing cholesterol values to be a function of age.

```
formula_2 <-  
  bf(bmi | mi() ~ age + mi(chl)) + # BMI model  
  bf(chl | mi() ~ age) # chl model  
m2 <- brm(formula_2,  
           data = nhanes)
```


# The output

```
summary(m2)
```

```
## Family: MV(gaussian, gaussian)
## Links: mu = identity; sigma = identity
##      mu = identity; sigma = identity
## Formula: bmi | mi() ~ age + mi(chl)
##      chl | mi() ~ age
## Data: nhanes (Number of observations: 25)
## Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
##      total post-warmup samples = 4000
##
## Population-Level Effects:
##      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## bmi_Intercept    21.94      3.82   14.65   29.84 1.00    1760    2101
## chl_Intercept    138.72     24.74   89.72  186.90 1.00    2654    2774
## bmi_age           -4.29      1.48   -7.17   -1.34 1.00    1274    2240
## chl_age           30.84     12.90    5.73   56.47 1.00    2707    2650
## bmi_michl         0.06      0.02    0.01    0.11 1.00    1070    1526
##
## Family Specific Parameters:
##      Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma_bmi       3.31      0.77    2.15    5.08 1.00    1558    1909
## sigma_chl       40.18      7.65   28.27   57.83 1.00    2095    2085
##
## Samples were drawn using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

## Visualizing the imputation

```
post<-posterior_samples(m2)  
plot(density(post$Ymi_bmi[10]))
```





**Pros** - We can allow for very complex structures in our models - We can neatly specify variable specific models for missing data - Sampling missing values using HMC as part of the model is consistent with Bayesian modeling principals

**Cons** - Gets technically complex quickly - Doesn't have easy solution for categorical data - Computationally intensive for big data

## Multiple imputation by chained equations (MICE)

MICE is a common pseudo-Bayesian approach to missing data. It produces  $k$  predictions for each unobserved variable based on a fully conditional regression models (where each variable is a function of all others).

- We'll allow each variable in the model to be a function of all others (we can relax this)
- Then make conditional predictions for each missing value
- This results in  $k$  complete datasets
- We apply our analysis over each dataset, and pool results

For a detailed description of the method and software, See Van Buuren and Groothuis-Oudshoorn, 2011:


<http://www.jstatsoft.org/v45/i03/paper>

It's honestly too easy

```
library(mice)
imp_nhanes<-mice(nhanes, m = 10, maxit = 20)
summary(imp_nhanes)
```


# Visualizing convergence

```
plot(imp_nhanes)
```



# Visualizing the imputations

```
stripplot(imp_nhanes, chl~.imp, pch=20, cex=2)
```



# Changing predictors

If we've got a collinear variable or a factor with many categories, we may want to exclude it from the model. Columns indicate predictors, row indicates outcomes.

```
pred<-imp_nhanes$predictorMatrix
```

Turn off hyp as a predictor of chl

```
pred["chl", "hyp"]<-0  
pred
```

```
##      age bmi hyp chl  
## age   0   1   1   1  
## bmi   1   0   1   1  
## hyp   1   1   0   1  
## chl   1   1   0   0
```

## Use the new predictor matrix

```
imp_nhanes2<-mice(nhanes, m = 10, maxit = 20,  
                  predictorMatrix = pred)  
summary(imp_nhanes)
```

## Change imputation methods

- The default in mice for continuous measures is partial mean matching
- The algorithm constructs a regression model, samples new coefficients from the parameter distribution, then makes a prediction. It then randomly selects one proximate observation from the observed data
- This makes the imputation data appear similar to the observed
- But sometimes we may not want to do this, or may have categorical data



# Imputation methods

For a full list of methods, see

<https://stefvanbuuren.name/fimd/sec-modelform.html>

```
meth<-imp_nhanes$method
```

Let's change chl to a linear model


```
meth[4]<-"norm"
```

And re-run the model

```
imp_nhanes3<-mice(nhanes, m = 10, maxit = 20,  
                  method = meth)  
summary(imp_nhanes3)
```

# Visualize imputations

```
densityplot(imp_nhanes3)
```



## Using the imputed data

After imputation, we conduct our analysis over each imputed dataset.

Typically, we'll need to separately fit each model.

For frequentist methods, either use a loop or the `with()` function

```
fit <- with(imp_nhanes3, lm(bmi ~ chl))
```

Then pool your results according to Rubin's rules for combination. This averages beta parameters, and adjusts standard errors for cross-imputation variance

# The results

```
pool(fit)
```

```
## Class: mipo    m = 10
##      term m   estimate      ubar      b      t dfcom
## 1 (Intercept) 10 21.31695134 1.113277e+01 2.129882e+00 1.347564e+01 23
## 2      chl 10  0.02587948 2.925939e-04 5.994597e-05 3.585345e-04 23
##      df      riv  lambda      fmi
## 1 16.56385 0.2104482 0.1738597 0.2583155
## 2 16.26681 0.2253655 0.1839170 0.2686309
```

In a Bayesian context, we don't have to worry about formulas for combination. We fit the model to each dataset, then pool the posterior samples from each for inference.

brms provides `brm_multiple()` to do this

```
brm_mi <- brm_multiple(bmi ~ age + chl,  
  data = imp_nhanes3, cores = 4)
```

```
summary(brm_mi)
```

```
## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: bmi ~ age + chl
## Data: imp_nhanes3 (Number of observations: 25)
## Samples: 40 chains, each with iter = 2000; warmup = 1000; thin = 1;
##           total post-warmup samples = 40000
##
## Population-Level Effects:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept      20.85      3.18   14.26   26.84 1.18      150      634
## age             -3.61      1.01   -5.59   -1.60 1.10      235     1224
## chl              0.06      0.02    0.02    0.10 1.19      143      685
##
## Family Specific Parameters:
##           Estimate Est.Error l-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma          3.08      0.54    2.21    4.33 1.09      261      713
##
## Samples were drawn using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

## The posterior samples

```
post<-posterior_samples(brm_mi)  
nrow(post)
```

```
## [1] 40000
```

- Measurement error is an important part of the data generating process
- Ignoring it presumes we have perfect measurement, can bias inference, and understates uncertainty
- Missing data is also important. List-wise deletion at best discards information (and overstates certainty), at worst biases inference
- If data is missing completely at random, or is conditional on an observed variable, we can impute the missings and recover valid inference
- If missingness is correlated with the outcome, we've got a problem



- Further reading on mice:  
<http://www.gerkovink.com/miceVignettes/>
- Further reading on Bayesian imputation and brms:  
[https://cran.r-project.org/web/packages/brms/vignettes/brms\\_missings.html](https://cran.r-project.org/web/packages/brms/vignettes/brms_missings.html)

# Thank you!

Thanks for a wonderful semester and for being experimental subjects as I try teaching Bayes!

Stay safe, and please be in touch: `frank.edwards@rutgers.edu`